

Logisland
event-mining@scale

@bailetthomas

Schedule

- Introduction
- Core concepts
- Knowledge Paradigm
- API Design
- Quick start

Introduction

Logisland Big picture

- **Multi-Purpose** realtime analytics framework
- **High scalability** and **Fault-tolerant**.
- **High throughput** (billions messages / day).
- **Easy** to operate on Hadoop or on **standalone containers**
- **Easily Extensible** to build high level apps
- **Open source**, initiated by Hurence

Use cases

- **Log aggregation** : low latency log processing
- **Stream processing** : multiple stages of processing (enriching, ...)
- **Complex Event processing** : write custom business Rules to generate alerts, for fraud detection
- **click stream tracking** : capture user click stream data
- **SIEM** : security manager for intrusion detection
- **IoT** : generate alerts based on outliers and forecasting.

Challengers

- **ELK** is great to start with, but hard to centralize processing and lacks of real offline ML
- **Splunk** is fantastic but clients are not rich enough to afford it ;)
- **NIFI** is a great tool but doesn't play well with distributed processing
- **Metron, Eagle** are security centric

Features

- **out-of-the-box components** (no code required)
- high level **extensible** framework
- raw data to structured records automatic **conversion**
- alert percolation or **query matching**
- event **governance** with Avro schema management
- **online prediction** with offline trained ML models

Features 2

- I/O to Elasticsearch, HBase, HDFS, RocksDB, ...
- telemetry sources (bro, pcap, netflow)
- live enrichment (geoip, custom lookups)
- SQL aggregations
- Time series **sampling**
- Outliers detection
- **Network footprint** clustering

Core Concepts

event =

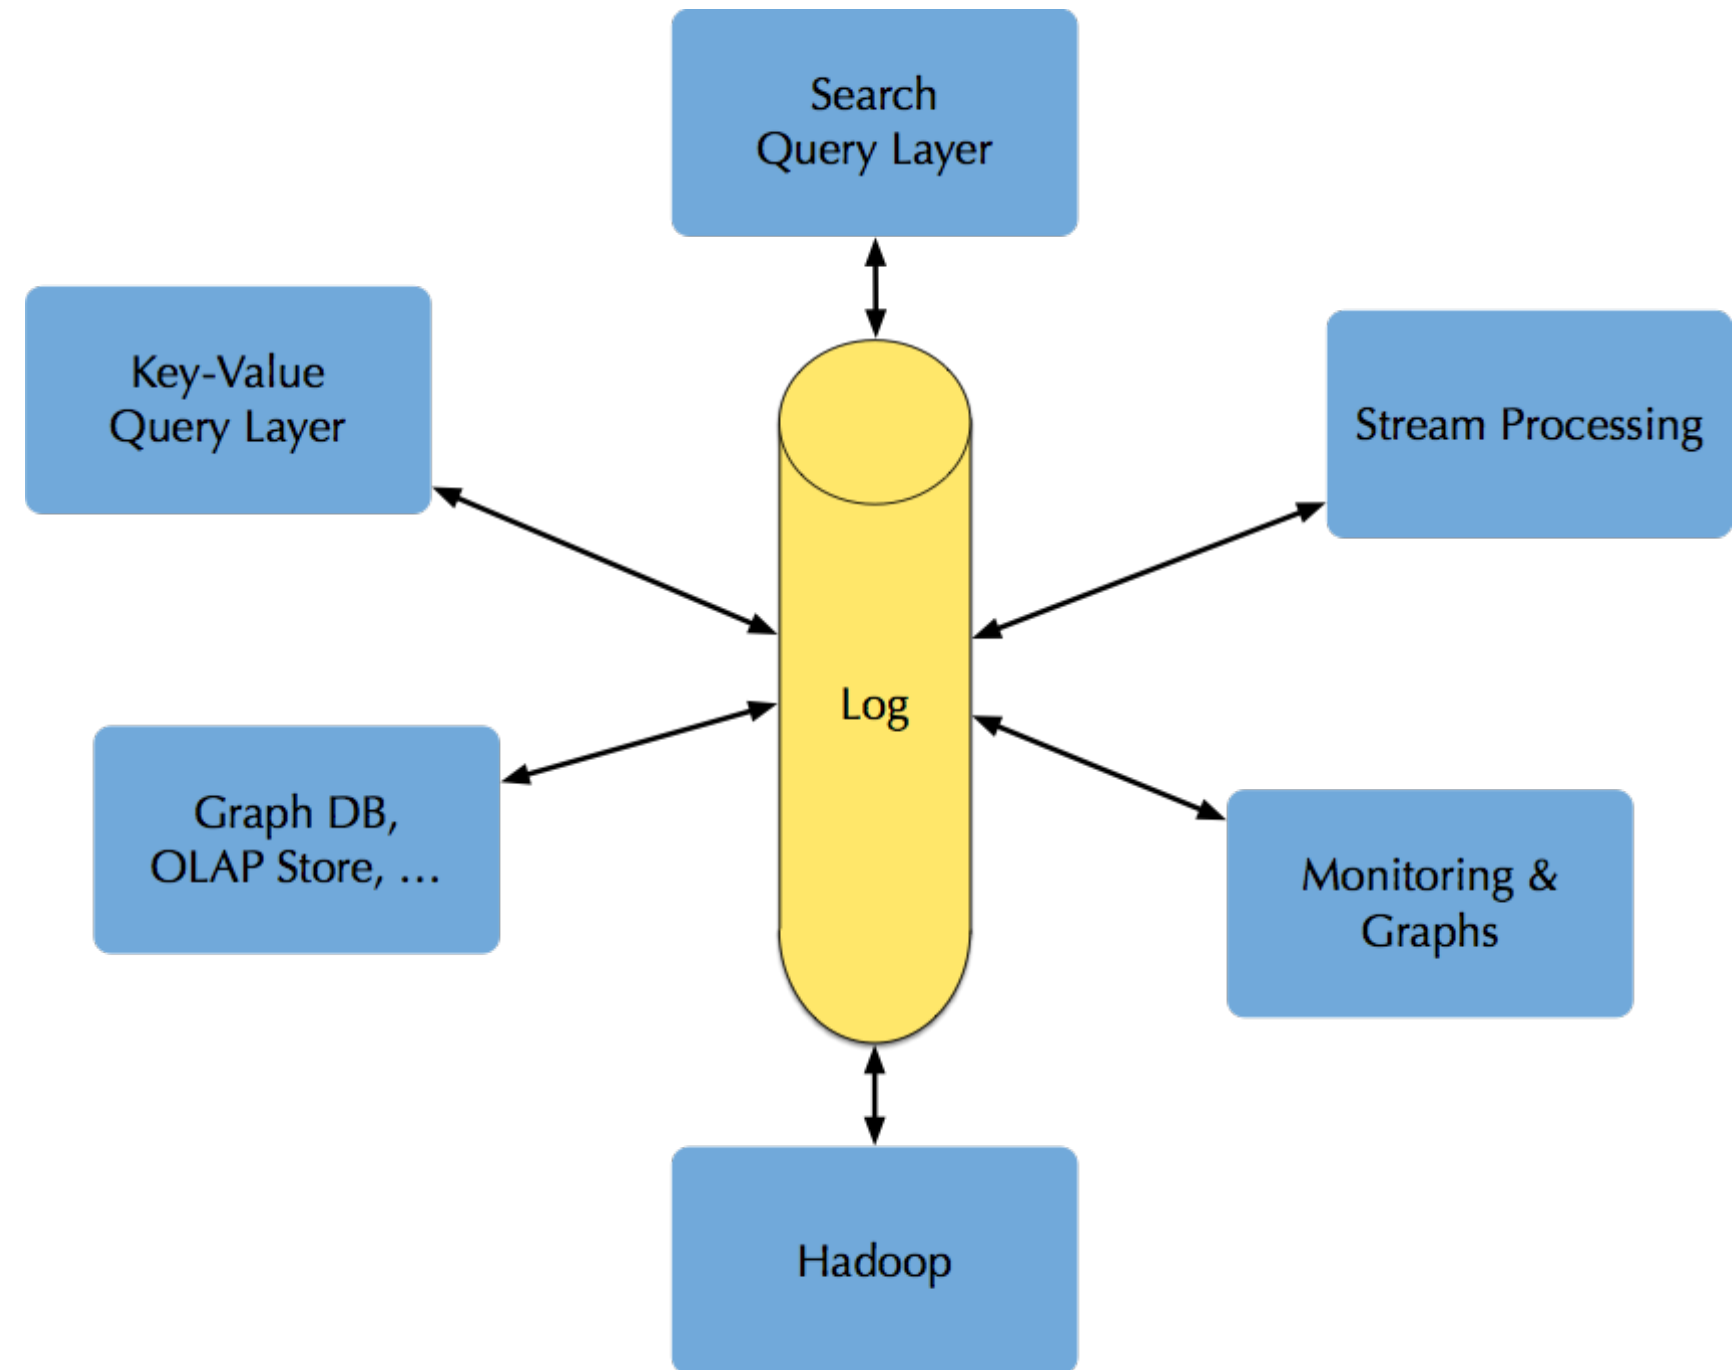
chronological change
in the system state

log =

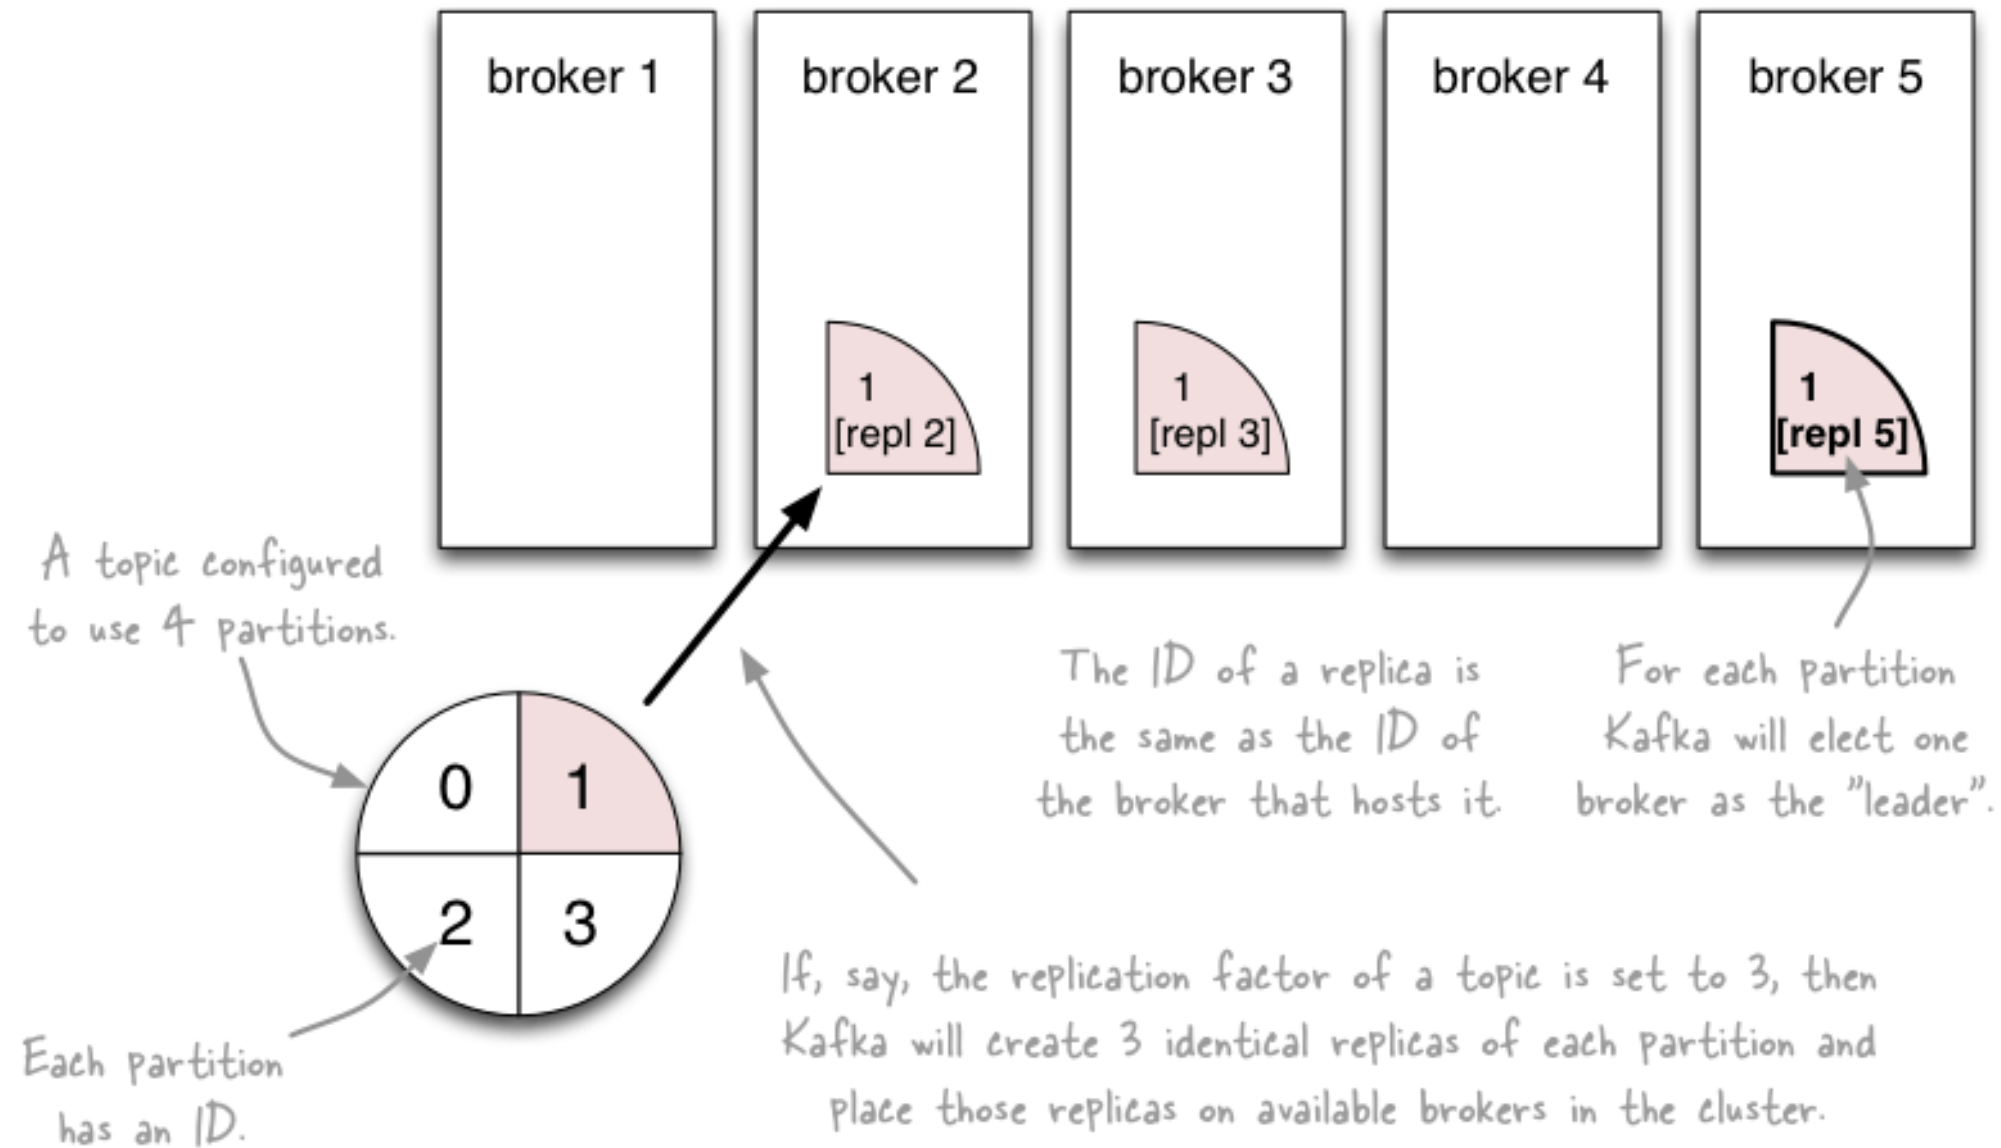
centralized registry of
chronologically ordered events

Log centric architecture

- **async** event production and consumption.
- **uncorrelated** publishers and subscribers.
- acts as a **Messaging system**.
- replay the log from any point in time.
- **realtime** event availability.



How to handle distributed logs ?





Logisland =

high level stream analytics solution
to handle massive scale
event processing

Pyramidal Knowledge Paradigm

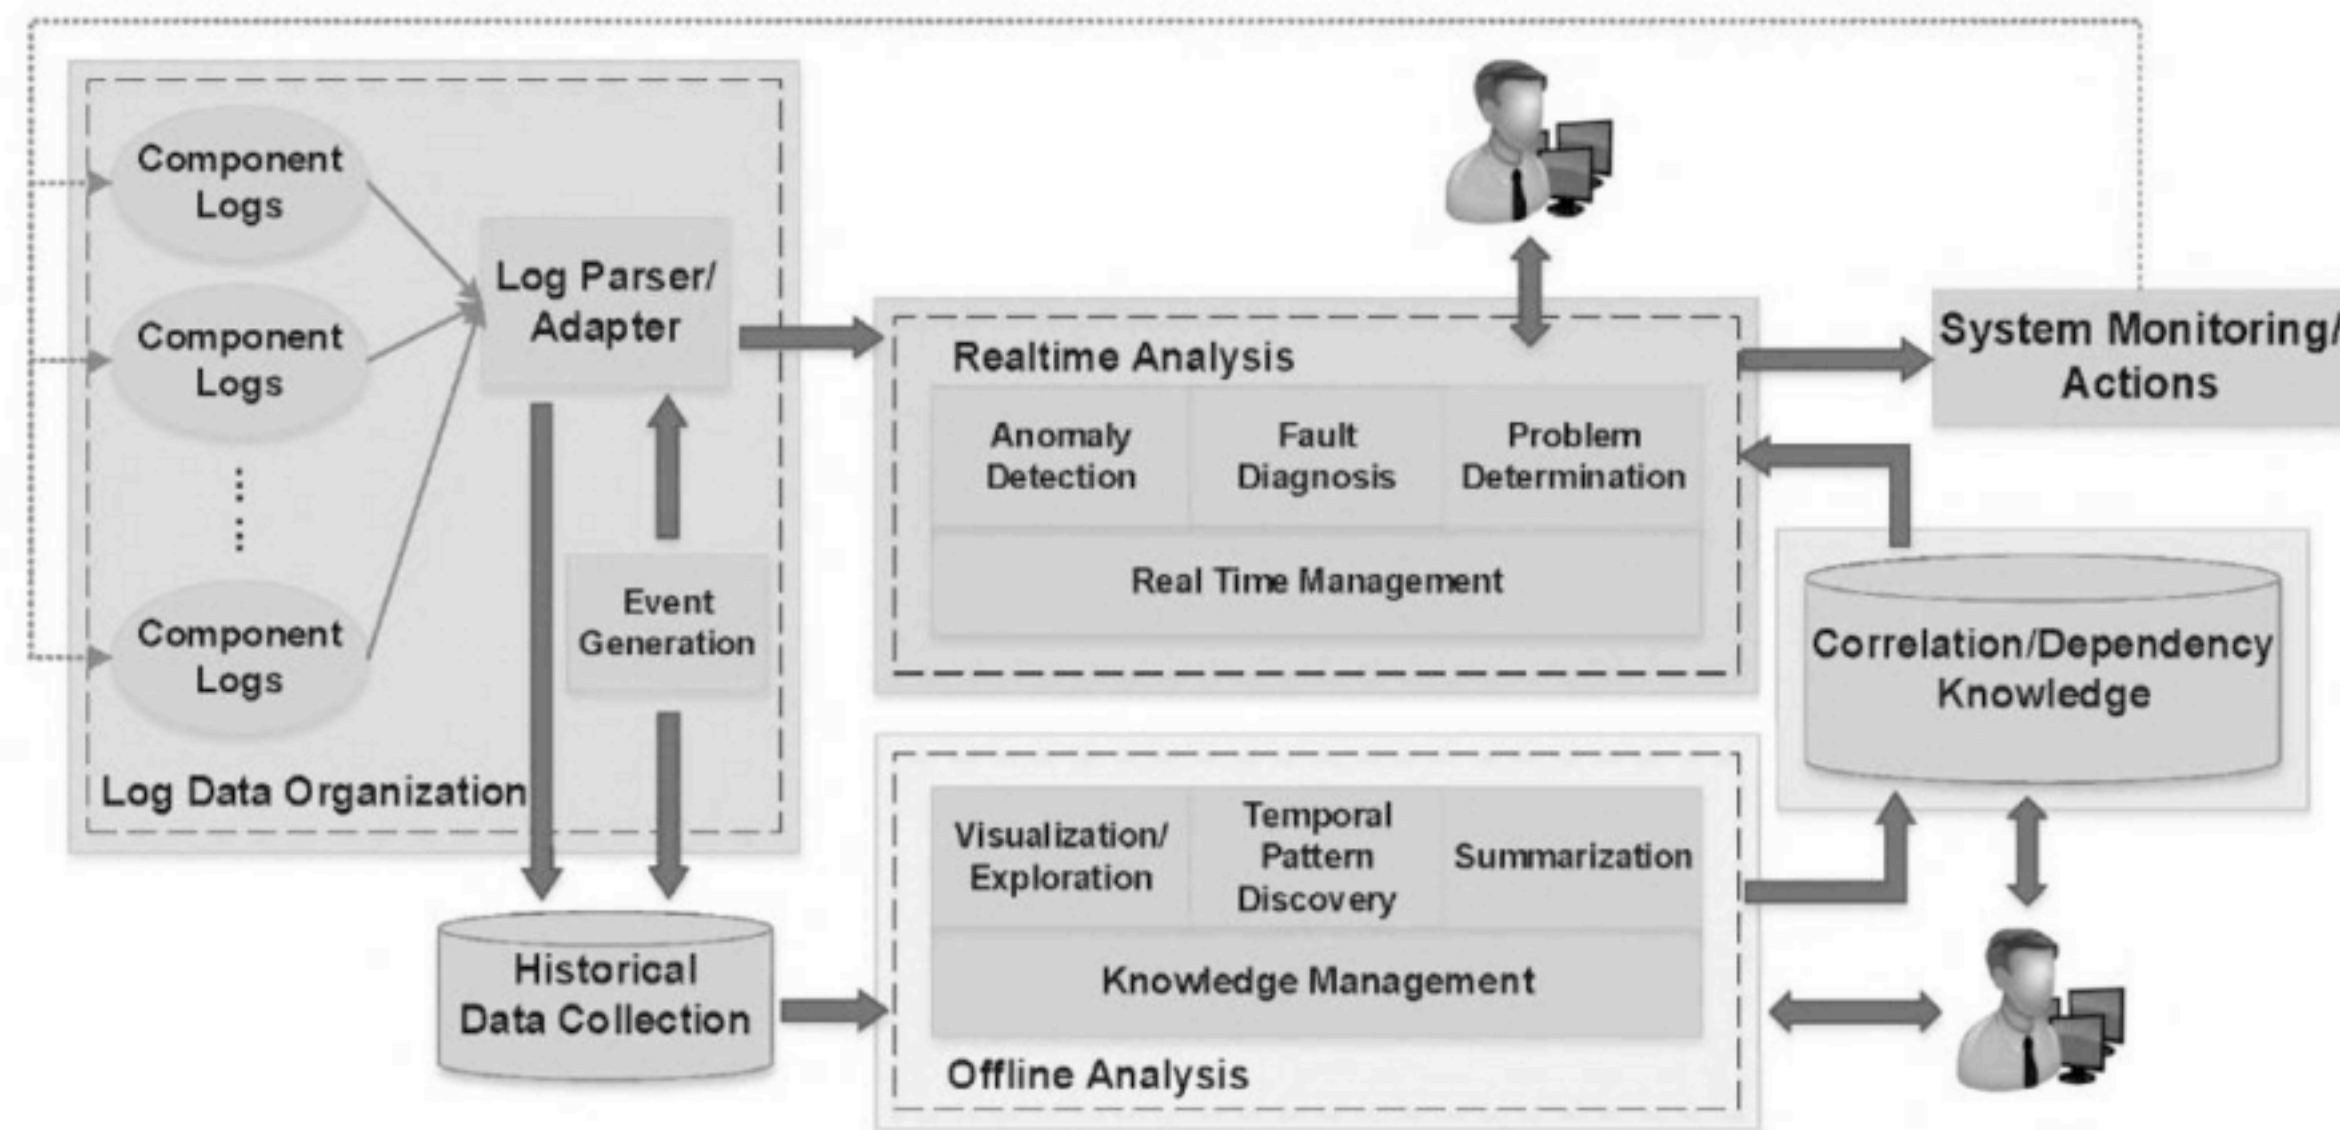
Logisland continuously transforms
data into information &
information into knowledge
by using asynchronous processing on
increasingly abstract
and meaningful records.





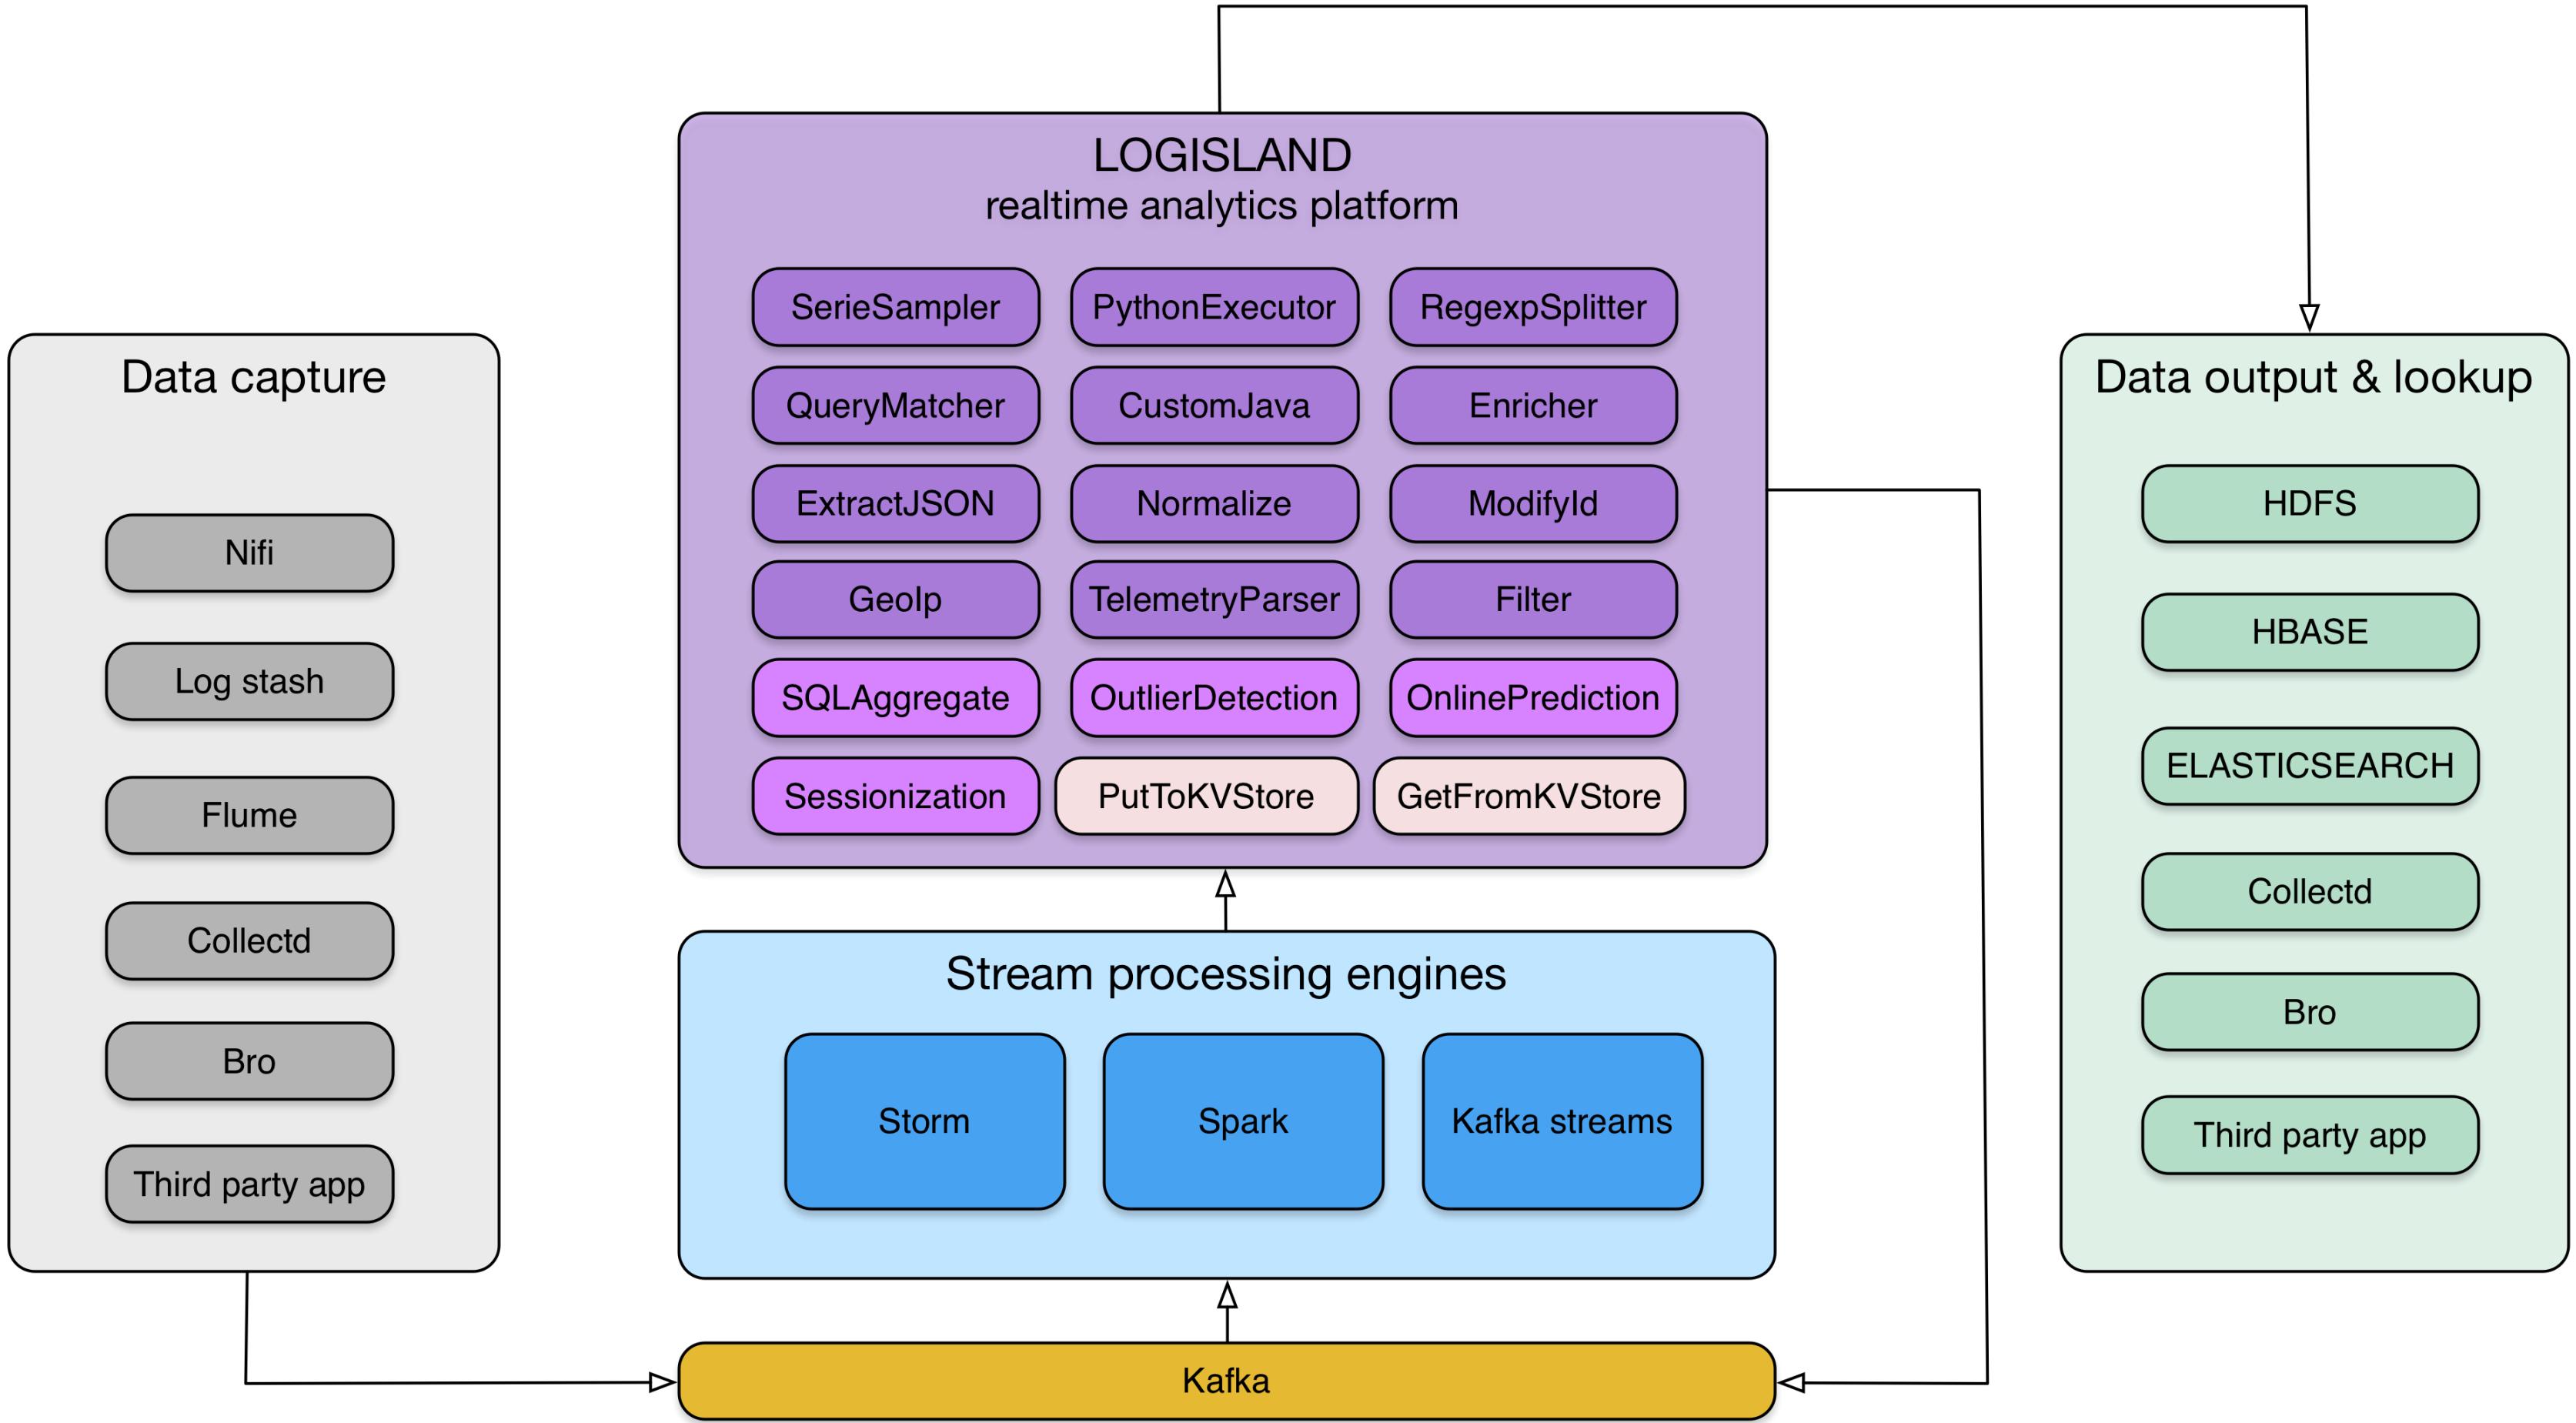
(credits : David McCandless, Information is Beautiful)

Data Driven Computing



API

Design



Record

The basic unit of processing is the Record.

A Record is a collection of Field, while a Field has a name, a type and a value.

```
String id = "firewall_record1";  
String type = "cisco";  
Record record = new Record(type).setId(id);  
  
assertTrue(record.isEmpty());  
assertEquals(record.size(), 0);
```

Field

A record holds a collection of fields.

```
record.setStringField("url_host", "origin-www.20minutes.fr")
    .setField("method", FieldType.STRING, "GET")
    .setField("response_size", FieldType.INT, 452)
    .setField("is_outside_office_hours", FieldType.BOOLEAN, false)
    .setField("tags",
        FieldType.ARRAY,
        Arrays.asList("spam", "filter", "mail"));

assertEquals( record.getField("method").asString(), "GET");
assertTrue( record.getField("response_size").asInteger() - 452 == 0);
record.removeField("is_outside_office_hours");
assertFalse( record.hasField("is_outside_office_hours"));
```

Special Field

A Record also has some special fields (type, time and id).

```
// shortcut for id
assertEquals(record.getId(), id);
assertEquals(record.getField(FieldDictionary.RECORD_ID).asString(),
              id);
```

```
// shortcut for time
assertEquals(record.getTime(),
              record.getField(FieldDictionary.RECORD_TIME).asLong());
```

```
// shortcut for type
assertEquals(record.getType(), type);
```

Field typing and validation

Fields are strongly typed, you can validate them

```
Record record = new StandardRecord();  
record.setField("request_size", FieldType.INT, 1399);  
assertTrue(record.isValid());
```

```
record.setField("request_size", FieldType.INT, "tom");  
assertFalse(record.isValid());
```

```
record.setField("request_size", FieldType.DOUBLE, 45.5d);  
assertTrue(record.isValid());
```

```
record.setField("request_size", FieldType.STRING, 45L);  
assertFalse(record.isValid());
```


Processor

Logisland is a **component centric** framework,

It's built over an abstraction layer to build **configurable components**.

A component can be Configurable and Configured.

The most common component you'll use is the Processor which takes a collection of Record and publish another collection of records

A configurable component that process Records

```
public interface Processor extends ConfigurableComponent {  
    void init(final ProcessContext context);  
  
    Collection<Record> process(ProcessContext context,  
                               Collection<Record> records);  
}
```

SplitText implementation

DefinePropertyDescriptor to handle components config.

```
@Tags({"parser", "regex", "log", "record"})
@CapabilityDescription("This is a processor that is used ...")
@dynamicProperty(name = "alternative regex & mapping", ...)
public class SplitText extends AbstractProcessor {

    public static final PropertyDescriptor VALUE_REGEX =
        new PropertyDescriptor.Builder()
            .name("value.regex")
            .description("the regex to match for the message value")
            .required(true)
            .addValidator(StandardValidators.NON_EMPTY_VALIDATOR)
            .build();

    ...
}
```

SplitText config

Use the components with simple yaml blocs.

```
- processor: apache_parser
  component: com.hurence.logisland.processor.SplitText
  type: parser
  documentation: a parser for apache log REGEX
  configuration:
    record.type: apache_log
    value.regex: (\S+)\s+(\S+)\s+(\S+)\s+\[([\w:\/] ...
    value.fields: src_ip,identd,user,record_time,http_method, ...
```

Stream

a record Stream basically :

- is a configurable Component
- reads a distributed collection of Record from Kafka input topics
- transmits them to a chain of Processor
- write the output collection of Record to some Kafka output topics

Streaming paradigm

You can handle partitioned data in 2 ways :

- **fully in parallel**, eg. a thread by partition, like with `KafkaRecordStreamParallelProcessing`, when records have no link with each other
- by **joining partitions** like with `KafkaRecordStreamSQLAggregator` or `KafkaRecordStreamHDFSBurner` when you need to join related records (costly join and shuffling operations)

Sample Stream configuration

Define a processing pipeline

```
- stream: parsing_stream
  component: com.hurence.logisland.stream.spark.KafkaRecordStreamParallelProcessing
  type: stream
  documentation: a processor that links
  configuration:
    kafka.input.topics: logisland_raw
    kafka.output.topics: logisland_events
    kafka.error.topics: logisland_errors
    kafka.input.topics.serializer: none
    kafka.output.topics.serializer: com.hurence.logisland.serializer.KryoSerializer
    kafka.error.topics.serializer: com.hurence.logisland.serializer.JsonSerializer
    ...
  processorConfigurations:
```

Engine

- The Engine manage a collection of Stream
- this is the abstraction of the execution model, mainly in Spark actually but plans are to integrate Beam to move on Storm and Kafka Streams
- you configure here your Spark job parameters

Sample engine configuration

Define a processing job

```
engine:
  component: com.hurence.logisland.engine.spark.KafkaStreamProcessingEngine
  type: engine
  documentation: Index some apache logs with logisland
  configuration:
    spark.app.name: IndexApacheLogsDemo
    spark.master: yarn-cluster
    spark.driver.memory: 1G
    spark.driver.cores: 1
    spark.executor.memory: 2G
    spark.executor.instances: 4
    spark.executor.cores: 2
    spark.yarn.queue: default
    ...
  streamConfigurations:
```

Transverse service injection : ControllerService

we often need to share access to external Services across the Processors, for example

- bulk buffers or client connections to external data
- a cache service that could cache K/V tuple across the worker node.

Sample ControllerService component

We need to provide an interface API for this service :

```
public interface CacheService<K,V> extends ControllerService {  
  
   PropertyDescriptor CACHE_SIZE = new PropertyDescriptor.Builder()  
        .name("cache.size")  
        .description("The maximum number of element in the cache.")  
        .required(false)  
        .defaultValue("16384")  
        .addValidator(StandardValidators.POSITIVE_INTEGER_VALIDATOR)  
        .build();  
  
    public V get(K k);  
  
    public void set(K k, V v);  
  
}
```

Inject service in Processor

You can then use this service in a custom processor :

```
public class TestProcessor extends AbstractProcessor {  
  
    static final PropertyDescriptor CACHE_SERVICE = new PropertyDescriptor.Builder()  
        .name("cache.service")  
        .description("CacheService")  
        .identifiesControllerService(CacheService.class)  
        .required(true)  
        .build();  
  
    @Override  
    public boolean hasControllerService() {  
        return true;  
    }  
}
```

Define service in config

The injection is done through yaml config files by injecting the instance of `lru_cache` Service.

```
controllerServiceConfigurations:
```

- controllerService: lru_cache
 component: com.hurence.logisland.service.elasticsearch.LRUKeyValueCacheService
 configuration:
 cache.size: 5000

```
streamConfigurations:
```

- stream: parsing_stream

```
processorConfigurations:
```

- processor: mock_processor
 component: com.hurence.logisland.processor.TestProcessorhing
 configuration:
 cache.service: lru_cache

Quick

Start

Getting started (Hadoop cluster)

Download the latest release from github

```
tar -xzf logisland-0.10.1-bin.tar.gz
```

Create a job configuration

```
vim conf/index-apache-logs.yml
```

Run the job

```
export SPARK_HOME=/usr/hdp/current/spark-client  
bin/logisland.sh --conf conf/index-apache-logs.yml
```

Getting started (lightweight container)

Pull & run the image from Docker Repository

```
docker pull hurence/logisland  
docker run -it --name logisland \  
    -p 8080:8080 -p 5601:5601 -p 9200:9200 \  
    -h sandbox hurence/logisland bash
```

Run the job

```
bin/logisland.sh --conf conf/index-apache-logs.yml
```


Next?

Roadmap

- Ambari Agent for job dynamic interaction (REST Api)
- visual Stream configuration / dashboards through Ambari views
- Auto-scaling to optimize cluster resources
- Density based automatic Usage profiling
- Pattern discovery through Deep Learning
- App store, per use-case knowledge bundles (cybersecurity, fraud, ...)

Resources

- **source** : <https://github.com/Hurence/logisland/releases>
- **Docker** : <https://hub.docker.com/r/hurence/logisland/tags/>
- **Documentation** : <http://logisland.readthedocs.io/en/latest/concepts.html>
- **support** : <https://gitter.im/logisland/logisland>
- **contact** : bailet.thomas@gmail.com

Q&A